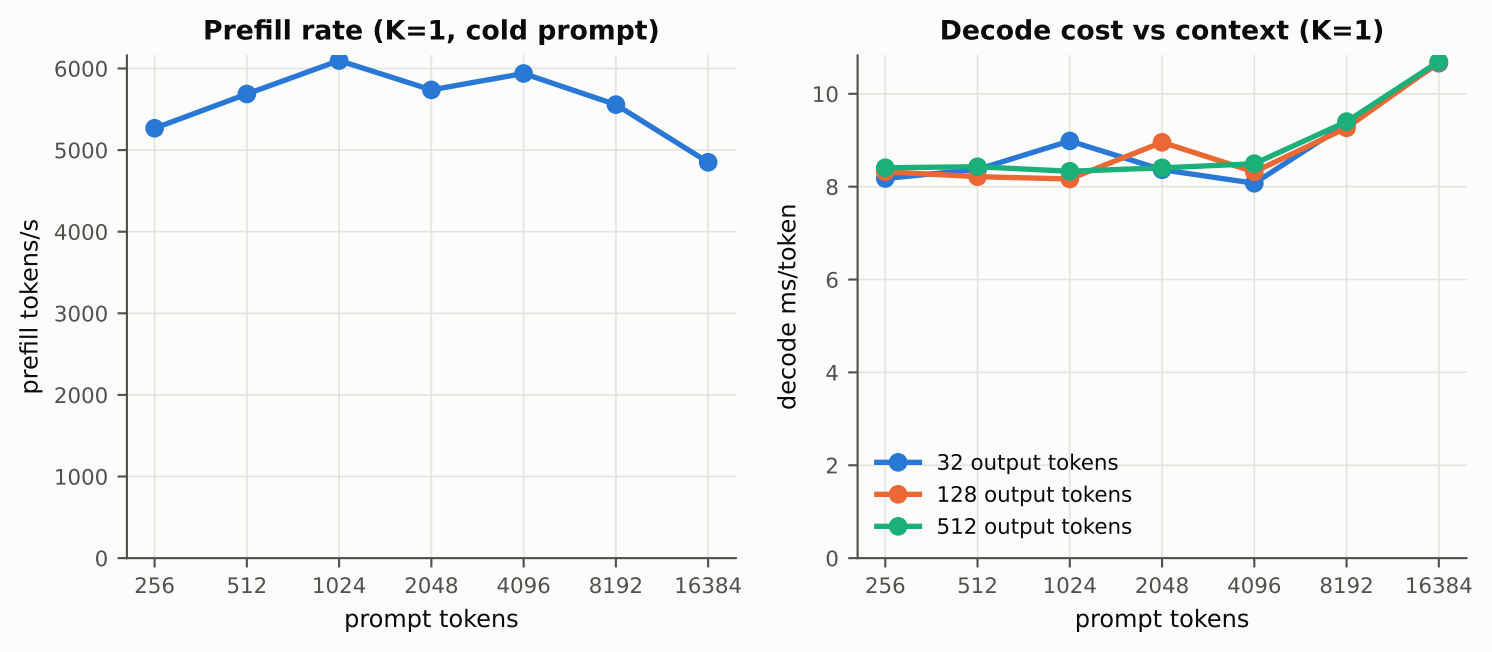
Prefill rate (K=1, cold prompt)
prefill tokens/s
0
1000
2000
3000
4000
5000
6000
256
512
1024
2048
4096
8192
16384
prompt tokens
Decode cost vs context (K=1)
decode ms/token
0
2
4
6
8
10
32 output tokens
128 output tokens
512 output tokens
256
512
1024
2048
4096
8192
16384
prompt tokens